

Task 2: Ensemble classifiers

Roger Llorenç, Victor Villegas, Luis Sierra

May 19, 2024

Contents

1	Introduction	1
2	Exploratory Data Analysis	1
3	Train-Test Split	7
4	Model fitting	7
4.1	Random Forest Classifier	8
4.2	Gradient Boosting	10
4.3	XGBoost	12
5	Predictor comparison	14

1 Introduction

For this assignment we be will employing predictive models based on decision trees to examine a dataset of cancer patients with the presence of 102 different peptides in the patients' organisms. The data is found in the files `cancerData.csv` and `cancerInfo.csv`.

```
[1]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import sklearn
from sklearn.metrics import accuracy_score, zero_one_loss
from sklearn.metrics import confusion_matrix, classification_report, ConfusionMatrixDisplay
from sklearn.model_selection import GridSearchCV
```

2 Exploratory Data Analysis

We first import the data and analyse its structure. We can immediately see that the data ingestion pipeline will require some work.

```
[2]: df1 = pd.read_csv("cancerDat.csv", sep = ";", decimal = ",")
df1 = df1.T
df1 = df1.drop('Unnamed: 0', axis=0)
```

```
display(df1.head())
print(df1.shape)
```

	0	1	2	3	4	5	\
NO.REC_1	21.923472	24.442617	19.050562	18.482667	24.086793	20.32946	
NO.REC_2	21.020165	23.649841	18.402413	19.088996	24.710323	21.495392	
NO.REC_3	19.585788	23.736128	18.191527	16.33124	21.917326	20.284533	
NO.REC_4	19.061767	23.374865	17.692775	15.36272	21.484924	18.379603	
NO.REC_5	18.547029	23.039588	19.066973	15.835721	21.339587	19.550809	

	6	7	8	9	...	92	\
NO.REC_1	19.304363	24.270429	18.878984	18.752264	...	19.439382	
NO.REC_2	19.454826	25.807051	19.091796	19.213397	...	20.631064	
NO.REC_3	16.853825	22.661125	18.215654	20.821777	...	19.123832	
NO.REC_4	16.513507	21.401436	18.38696	19.847221	...	17.958307	
NO.REC_5	16.831653	21.776832	17.85408	20.368534	...	18.212854	

	93	94	95	96	97	98	\
NO.REC_1	16.142102	22.858297	22.262118	18.079186	17.151515	20.912124	
NO.REC_2	NaN	22.028998	22.031468	17.101384	18.315637	21.512601	
NO.REC_3	16.171227	23.143305	22.334392	17.159968	16.859732	19.076147	
NO.REC_4	NaN	19.183961	17.851328	16.564709	15.161135	18.190653	
NO.REC_5	NaN	22.228449	21.385404	17.072001	15.071656	18.896095	

	99	100	101
NO.REC_1	17.298159	19.097263	21.21211
NO.REC_2	17.100711	18.875548	23.980238
NO.REC_3	15.417028	16.340283	19.810886
NO.REC_4	15.269443	15.411408	18.351433
NO.REC_5	15.147357	NaN	20.28779

[5 rows x 102 columns]

(129, 102)

```
[3]: df2 = pd.read_csv("cancerInfo.csv", sep = ";")
df2 = df2.drop(['Unnamed: 0', 'Group'], axis = 1)
df2.index = df2.iloc[:,0]
df2 = df2.drop(['sampleNames'], axis = 1)
display(df2.head())
```

	sites
sampleNames	
NO.REC_1	A
NO.REC_2	B
NO.REC_3	C
NO.REC_4	C
NO.REC_5	C

We want to define a single response variable according to whether the cancer sees a recurrence or not. This is currently coded into the key of the patient, so we will create a separate binary vector to encode this response, where 0 will denote non-recurrence and 1 will denote recurrence.

```
[4]: indices1 = np.array(df1.index).astype(str)
index_array1 = np.zeros(len(indices1))
labels_1 = np.where(np.char.startswith(indices1, 'NO.REC_'), 0, np.where(np.char.
    ↳startswith(indices1, 'REC_'), 1, index_array1))

indices2 = np.array(df2.index).astype(str)
index_array2 = np.zeros(len(indices2))
labels_2 = np.where(np.char.startswith(indices2, 'NO.REC_'), 0, np.where(np.char.
    ↳startswith(indices2, 'REC_'), 1, index_array2))

print("Labels are equal: " + str(np.array_equal(labels_1, labels_2))) # That is_
    ↳we can define a single response variable y for the two files
y = labels_1
```

Labels are equal: True

```
[5]: y = pd.Series(y)
display(y)
```

```
0      0.0
1      0.0
2      0.0
3      0.0
4      0.0
```

```
...
```

```
124     1.0
125     1.0
126     1.0
127     1.0
128     1.0
```

Length: 129, dtype: float64

Now we will check for missing values and some descriptive statistics of the data. There is a substantial amount of missing values in the dataset, so in order to avoid discarding such a large amount of information, we resort to imputing the linear interpolation of the missing entries. This increases the amount of information that we work with, but we claim that for the final results this is preferable to eliminating the entire rows and rendering our dataset almost empty.

```
[6]: for col in df1:
      df1[col] = pd.to_numeric(df1[col], errors='coerce')
df1.dtypes
```

```
[6]: 0      float64
      1      float64
```

```

2      float64
3      float64
4      float64
...
97     float64
98     float64
99     float64
100    float64
101    float64
Length: 102, dtype: object

```

```

[7]: print(df1.isnull().values.any())
      print(sum(df1.isnull()))
      df1 = df1.interpolate(method='linear', limit_direction='forward')
      print(df1.isnull().values.any())

```

```

True
5151
False

```

```

[8]: print(df2.isnull().values.any()) # False -> No nulls

```

```
False
```

Now that we have handled the missing values, we can go ahead and perform some more sophisticated exploration techniques.

```

[9]: df1.describe()

```

```

[9]:
count      0      1      2      3      4      5  \
mean    129.000000  129.000000  129.000000  129.000000  129.000000  129.000000
std       19.930039   23.286214   18.226328   17.331228   22.591467   20.193446
min       1.646143    0.784218    1.126192    1.477403    1.562914    1.185180
25%      15.836862   21.313713   14.401524   12.976690   18.807566   16.963778
50%      18.547029   22.755110   17.440432   16.166380   21.321595   19.397560
75%      20.341514   23.374865   18.349787   17.373647   22.461809   20.264099
max       21.362745   23.793517   18.968572   18.501081   23.992847   20.961033

count      6      7      8      9  ...    92  \
mean    129.000000  129.000000  129.000000  129.000000  ...  129.000000
std      18.133973   23.193300   18.687522   19.516615  ...   19.126770
min      1.751196    1.860021    1.744740    1.068425  ...    1.399915
25%     13.770822   18.824646   16.818849   16.952402  ...   14.488970
50%     17.001766   22.036426   18.109188   18.830477  ...   18.323381
75%     18.421874   23.469361   18.548956   19.546749  ...   19.240883
max     19.426085   24.633681   18.967436   20.108405  ...   19.956155

count      6      7      8      9  ...    92  \
mean    21.934433   26.596796   36.306672   22.920624  ...   23.668357
std      21.934433   26.596796   36.306672   22.920624  ...   23.668357
min      21.934433   26.596796   36.306672   22.920624  ...   23.668357
25%      21.934433   26.596796   36.306672   22.920624  ...   23.668357
50%      21.934433   26.596796   36.306672   22.920624  ...   23.668357
75%      21.934433   26.596796   36.306672   22.920624  ...   23.668357
max      21.934433   26.596796   36.306672   22.920624  ...   23.668357


```

	93	94	95	96	97	98 \
count	129.000000	129.000000	129.000000	129.000000	129.000000	129.000000
mean	15.832452	21.643638	21.312564	17.060424	17.096818	19.891371
std	1.005218	1.009982	1.126496	0.881310	1.241427	1.487108
min	12.398937	18.802740	17.851328	15.050081	13.746502	15.860695
25%	15.358616	20.940953	20.538005	16.466670	16.297205	18.855134
50%	16.099669	21.663512	21.281094	17.156167	17.270199	19.892329
75%	16.514396	22.295952	22.004275	17.734682	18.003294	21.065003
max	19.057885	24.195671	23.932574	19.008904	20.339064	22.552429

	99	100	101
count	129.000000	129.000000	129.000000
mean	16.372295	17.653626	20.692479
std	1.011885	1.123423	2.025807
min	13.846473	14.930870	15.644922
25%	15.707511	17.008565	19.477624
50%	16.416639	17.691291	20.655011
75%	17.101205	18.540819	22.009082
max	18.586913	19.888654	25.794838

[8 rows x 102 columns]


```
[10]: df1.hist(figsize=(12,10))
plt.tight_layout()
plt.show()
```



In the previous cells, we can observe some erratic distributions for some of the covariates, where they do not seem to fit very well with the statement of the law of large numbers. Specifically, as an example, the covariates indexed in positions 15, 54 or 78 have some interesting shapes.

```
[11]: from sklearn.decomposition import PCA
import itertools
pca = PCA(n_components=20)
pca.fit_transform(df1)
#print(pca.explained_variance_ratio_)
#print(list(itertools.accumulate(pca.explained_variance_ratio_)))
plt.plot(list(itertools.accumulate(pca.explained_variance_ratio_)))
```

```
[11]: [<matplotlib.lines.Line2D at 0x779719e76560>]
```



We can see from the PCA that the amount of variance accumulated by the twenty first principal components amounts to almost 90%. Our final classifiers will only include ten peptides as features to train on.

3 Train-Test Split

```
[12]: from sklearn.model_selection import train_test_split
      np.random.seed(1234)
```

We separate the data into two sets: the training set and the test set.

```
[13]: X_train, X_test, y_train, y_test = train_test_split(
      df1,
      y,
      test_size = 1/3,
      random_state = 1234)
```

```
[14]: print(X_train.shape, X_test.shape, y_train.shape, y_test.shape)
```

```
(86, 102) (43, 102) (86,) (43,)
```

4 Model fitting

We'll fit the following models: a Random Forest Classifier and two Boosting Classifiers.

4.1 Random Forest Classifier

```
[15]: from sklearn.ensemble import RandomForestClassifier
      rf = RandomForestClassifier(random_state = 1234)
      parameters = {'n_estimators' : range(10,130,20),
                    'max_features' : (1,2,3,4,5,6,7,8,9,10)
                    }
```

We implement a grid search procedure.

```
[16]: grid_rf = GridSearchCV(rf, parameters)
      grid_rf.fit(X_train, y_train)
```

```
[16]: GridSearchCV(estimator=RandomForestClassifier(random_state=1234),
                  param_grid={'max_features': (1, 2, 3, 4, 5, 6, 7, 8, 9, 10),
                              'n_estimators': range(10, 130, 20)})
```

```
[17]: rf_results = pd.DataFrame(grid_rf.cv_results_)
      rf_results.filter(regex = '(param.*|mean_t|std_t)') \
        .drop(columns = 'params') \
        .sort_values('mean_test_score', ascending = False) \
        .head(4)
```

```
[17]:   param_max_features  param_n_estimators  mean_test_score  std_test_score
36                   7                   10         0.594118         0.118234
55                  10                   30         0.582353         0.079792
48                   9                   10         0.572549         0.142693
1                    1                   30         0.569935         0.044213
```

```
[18]: best_params = grid_rf.best_params_
      print("Best Parameters:", best_params)
      best_rf = RandomForestClassifier(**best_params, random_state=1234)
      best_rf.fit(X_train, y_train)
```

Best Parameters: {'max_features': 7, 'n_estimators': 10}

```
[18]: RandomForestClassifier(max_features=7, n_estimators=10, random_state=1234)
```

The accuracy and performance of RF is as follows.

```
[19]: y_pred = best_rf.predict(X_test)
      accuracy_score(y_pred, y_test)
```

```
[19]: 0.5348837209302325
```


```
[20]: confusion_mtx = confusion_matrix(y_test, y_pred)
      disp = ConfusionMatrixDisplay(confusion_matrix= confusion_mtx, display_labels=
      ↪None)
```



```

disp.plot(cmap= plt.cm.Blues)
plt.title("Confusion matrix Random Forest")
plt.show()

```



Finally, the most relevant variables in the prediction are given by the following code:

```

[21]: feature_importances = best_rf.feature_importances_
importance_df = pd.DataFrame({'Feature' : X_train.columns, 'Importance' : ↵
↵feature_importances})
print(importance_df.sort_values(by=['Importance'], ascending=False).head(10))

```

	Feature	Importance
51	51	0.057151
11	11	0.043937
25	25	0.032852
96	96	0.030232
78	78	0.029687
76	76	0.028671
5	5	0.027653
34	34	0.026475
12	12	0.026456
62	62	0.024448

4.2 Gradient Boosting

Using stumps as classification trees for the response variable, we compute the misclassification rates of both the learning set and the test set across 2,000 iterations.

```
[22]: from sklearn.ensemble import GradientBoostingClassifier
gb = GradientBoostingClassifier(n_estimators=2000, random_state=1234)
parameters = {'max_depth': (1,4,8,16)}
```

```
[23]: grid_gb = GridSearchCV(gb, parameters)
grid_gb.fit(X_train, y_train)
```

```
[23]: GridSearchCV(estimator=GradientBoostingClassifier(n_estimators=2000,
                                                         random_state=1234),
                  param_grid={'max_depth': (1, 4, 8, 16)})
```

The results of comparing the test-set misclassification rates attained by different ensemble classifiers based on trees with varying maximum depth is shown in the following table.

```
[24]: gb_results = pd.DataFrame(grid_gb.cv_results_)
gb_results.filter(regex = '(param.*|mean_t|std_t)') \
    .drop(columns = 'params') \
    .sort_values('mean_test_score', ascending = False) \
    .head(4)
```

```
[24]:   param_max_depth  mean_test_score  std_test_score
0                1          0.524183         0.069834
2                8          0.454248         0.103714
3               16          0.454248         0.103714
1                4          0.442484         0.090319
```

The following plot represents graphically the error as a function of the number of boosting iterations.

```
[25]: best_params = grid_gb.best_params_
print("Best Parameters:", best_params)


train_errors = []
best_gb = GradientBoostingClassifier(**best_params, random_state=1234)
best_gb.fit(X_train, y_train)

for i, y_pred_train in enumerate(best_gb.staged_predict(X_train)):
    train_errors.append(zero_one_loss(y_train, y_pred_train))

plt.plot(np.arange(0,100), train_errors)
plt.xlabel('Number of Boosting iterations')
plt.ylabel('Misclassification rate')
plt.title('Gradient Boosting Error Evolution')
```

Best Parameters: {'max_depth': 1}

```
[25]: Text(0.5, 1.0, 'Gradient Boosting Error Evolution')
```

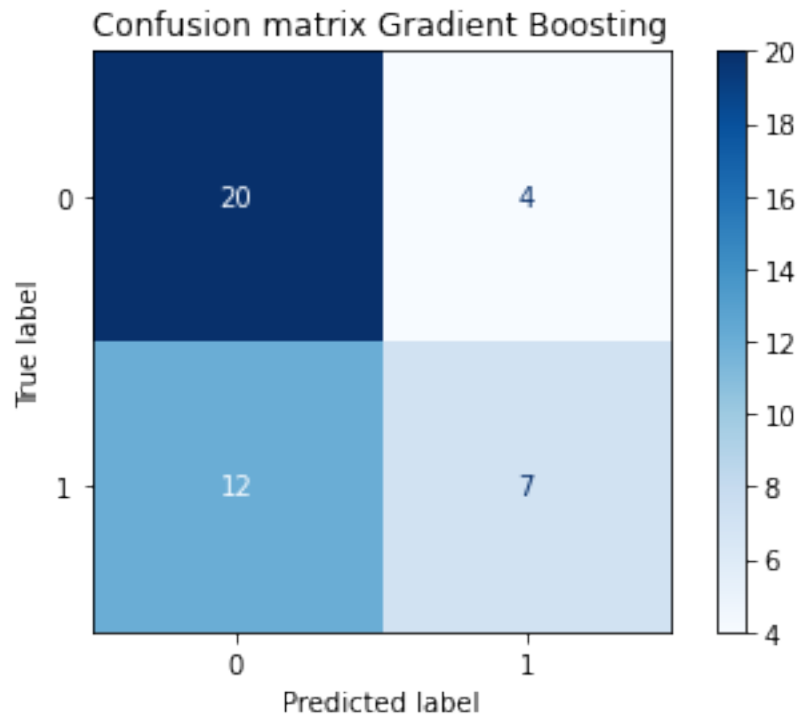


The accuracy and performance of Gradient Boosting is as follows.

```
[34]: y_pred = best_gb.predict(X_test)
accuracy_score(y_pred, y_test)
```

```
[34]: 0.627906976744186
```

```
[35]: confusion_mtx = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix= confusion_mtx, display_labels=
    ↪None)
disp.plot(cmap= plt.cm.Blues)
plt.title("Confusion matrix Gradient Boosting")
plt.show()
```



4.3 XGBoost

To conclude, we introduce another boosting flavour and analyse its performance.

```
[36]: import xgboost
xgb = xgboost.XGBClassifier(n_estimators=2000, random_state=1234)
parameters = {'max_depth' : (1,4,8,16),
              'learning_rate' : [0,2]
              }

[37]: grid_xgb = GridSearchCV(xgb, parameters)
grid_xgb.fit(X_train, y_train)
```

```
[37]: GridSearchCV(estimator=XGBClassifier(base_score=None, booster=None,
callbacks=None, colsample_bylevel=None,
colsample_bynode=None,
colsample_bytree=None, device=None,
early_stopping_rounds=None,
enable_categorical=False, eval_metric=None,
feature_types=None, gamma=None,
grow_policy=None, importance_type=None,
interaction_constraints=None,
learning_rate=None, max_bin=None,
max_cat_threshold=None,
```

```

max_cat_to_onehot=None,
max_delta_step=None, max_depth=None,
max_leaves=None, min_child_weight=None,
missing=nan, monotone_constraints=None,
multi_strategy=None, n_estimators=2000,
n_jobs=None, num_parallel_tree=None,
random_state=1234, ...),
param_grid={'learning_rate': [0, 2], 'max_depth': (1, 4, 8, 16)})

```

```

[38]: xgb_results = pd.DataFrame(grid_xgb.cv_results_)
xgb_results.filter(regex = '(param.*|mean_t|std_t)') \
      .drop(columns = 'params') \
      .sort_values('mean_test_score', ascending = False) \
      .head(4)

```

```

[38]:  param_learning_rate  param_max_depth  mean_test_score  std_test_score
0                0           1          0.581699          0.013072
1                0           4          0.581699          0.013072
2                0           8          0.581699          0.013072
3                0          16          0.581699          0.013072

```

```

[39]: best_params = grid_xgb.best_params_
print("Best Parameters:", best_params)

best_xgb = xgboost.XGBClassifier(**best_params, random_state=1234)
best_xgb.fit(X_train, y_train)

```

Best Parameters: {'learning_rate': 0, 'max_depth': 1}

```

[39]: XGBClassifier(base_score=None, booster=None, callbacks=None,
      colsample_bylevel=None, colsample_bynode=None,
      colsample_bytree=None, device=None, early_stopping_rounds=None,
      enable_categorical=False, eval_metric=None, feature_types=None,
      gamma=None, grow_policy=None, importance_type=None,
      interaction_constraints=None, learning_rate=0, max_bin=None,
      max_cat_threshold=None, max_cat_to_onehot=None,
      max_delta_step=None, max_depth=1, max_leaves=None,
      min_child_weight=None, missing=nan, monotone_constraints=None,
      multi_strategy=None, n_estimators=None, n_jobs=None,
      num_parallel_tree=None, random_state=1234, ...)

```

The accuracy and performance of XGBoost is as follows.

```

[40]: y_pred = best_xgb.predict(X_test)
accuracy_score(y_pred, y_test)


```

```

[40]: 0.5581395348837209

```

```
[41]: confusion_mtx = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix= confusion_mtx, display_labels=
    ↪None)
disp.plot(cmap= plt.cm.Blues)
plt.title("Confusion matrix XGBoost")
plt.show()
```



We can see how the gradient boosting always predicts the most common class, and so has a large amount of Type 1 errors.

5 Predictor comparison

The accuracies of the three proposed models are the following:

```
[45]: print("Accuracy of RF:", accuracy_score(best_rf.predict(X_test), y_test))
print("Accuracy of GB:", accuracy_score(best_gb.predict(X_test), y_test))
print("Accuracy of XGB:", accuracy_score(best_xgb.predict(X_test), y_test))
```

```
Accuracy of RF: 0.5348837209302325
Accuracy of GB: 0.627906976744186
Accuracy of XGB: 0.5581395348837209
```

We can see that, in general, they're not very high. This can be due to the number of NaNs in the data used. However, the Gradient Boosting method obtains a result considerably better than the

other two methods, followed by the XGB and then the RF, which obtain similar results. For this reason, the classifier we would choose would be Gradient Boosting.

Using Gradient Boosting, the ten most relevant recurrence biomarkers are the following peptides:

```
[46]: feature_importances = best_gb.feature_importances_  
importance_gb = pd.DataFrame({'Feature' : X_train.columns, 'Importance' :  
    ↪ feature_importances})  
print(importance_gb.sort_values(by=['Importance'], ascending=False).head(10))
```

	Feature	Importance
51	51	0.123774
11	11	0.112551
65	65	0.089577
60	60	0.067424
38	38	0.067323
70	70	0.067106
76	76	0.058383
19	19	0.048530
75	75	0.036543
67	67	0.032089